Performance is not (only) about micro-optimizations!

Clément Grégoire

MUREX

CODE RECKONS

Science to the CORE

Performance is not (only) about micro-optimizations!

Who am I?

Clément GREGOIRE

- 10 years of C++ Programming
- 5 years in game-dev
  - Microsoft Flight Simulator
  - Various ports
- Software optimization is my job
- Co-founder & consultant

siliceum

This talk is

- Not meant to convince you performance is important
- Not your usual micro-benchmarking talk (obviously)
- Not specific to C++

"Performance Matters" by Emery Berger
https://youtu.be/r-TLSBdHe1A

siliceum

Software performance

What's that?

Speed

How long does it take for a task to execute?
Latency?

- CPU
- GPU
- Network
- Startup / loading time

Memory usage

How much memory can we afford to use?

Do we need to play nice with other applications?

- Embedded
- Cloud
- Video games
  - PC vs Console
- Multi-tasking

Storage

What is my package size?
How long will it take to download?

- Embedded
- Docker images
- Games
  - Textures, sounds, 3D models, ...
- Web
  - Images, minification, ...

“Fallout 76’s Day One Patch is Larger Than the Actual Game”

Resilience

- Security (DoS)
- Scaling up
- Availability

Maintainability

- Speed
  - Development / prototyping
  - Identification of performance issues
- Iteration/build time (Live++ anyone?)
- Debug mode performance

siliceum

Optimizations

Where to begin?

The optimization cycle

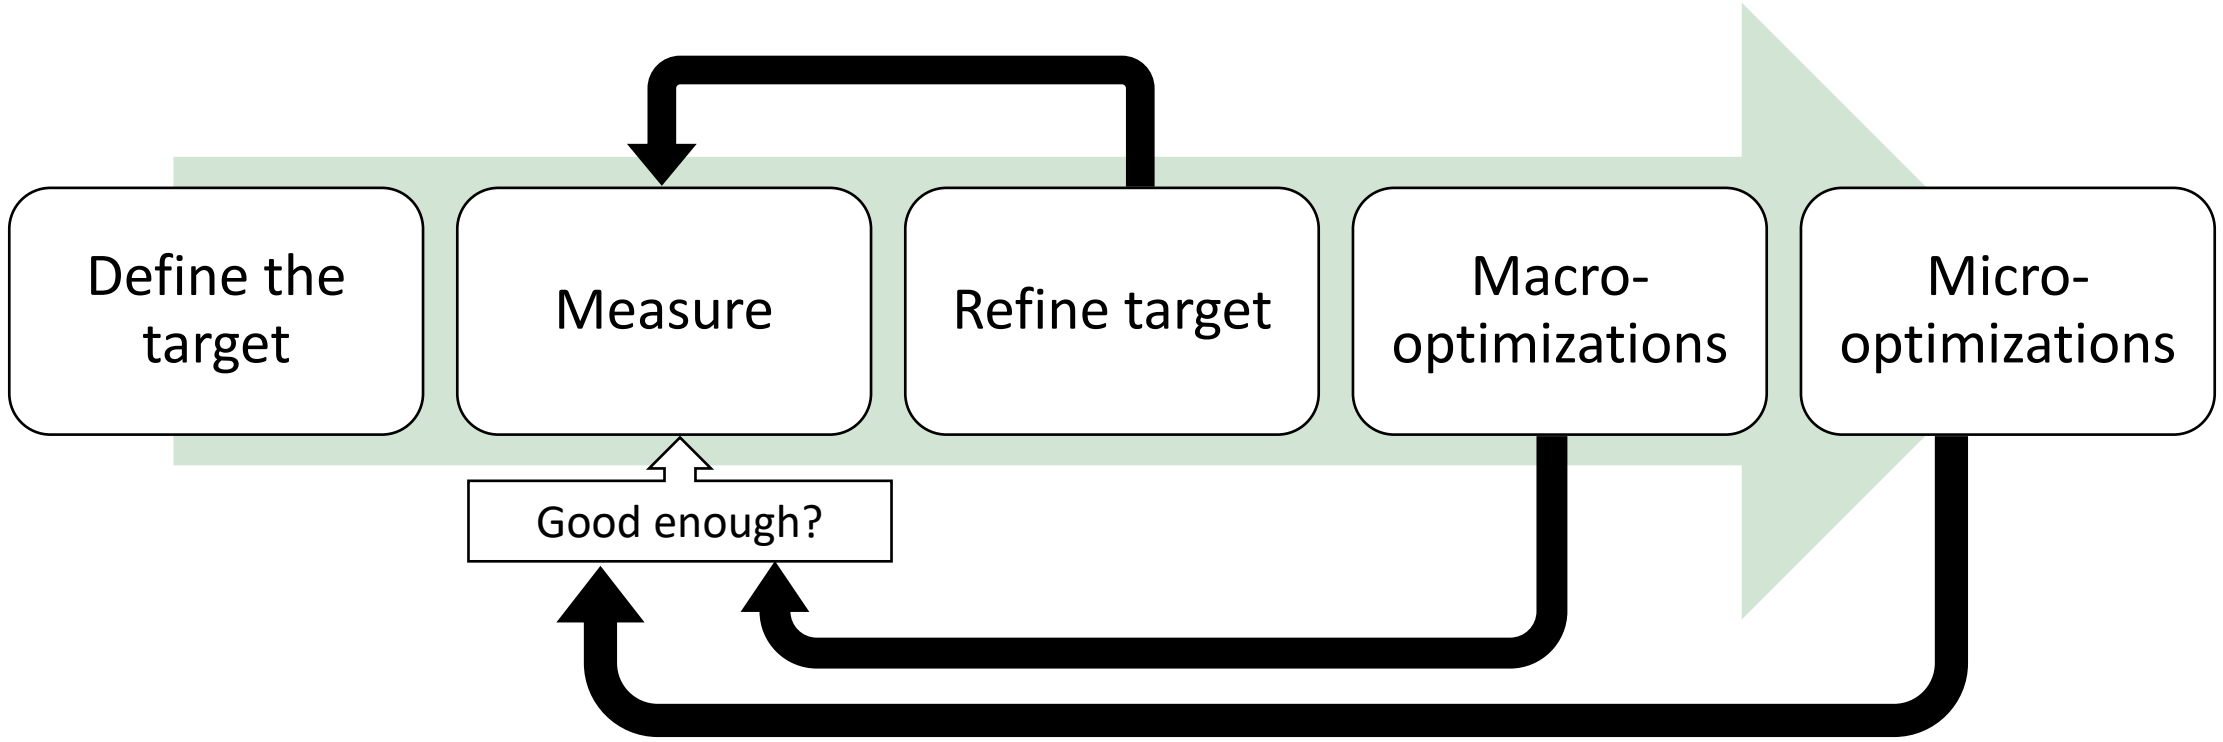

Defining performance objectives

A few examples

- 60 frames per second
- 8GB of memory maximum
- Web page loading < 1s
- Be the best (=> no target, only if you have time to afford it)

Defining performance objectives

Feasibility and cost-effectiveness

- Can this actually be done?
- How long will it take?
- Is it an isolated use-case?
- Buy better hardware?
  (yes but **no**, end of Moore's law, Amdahl's law, …)

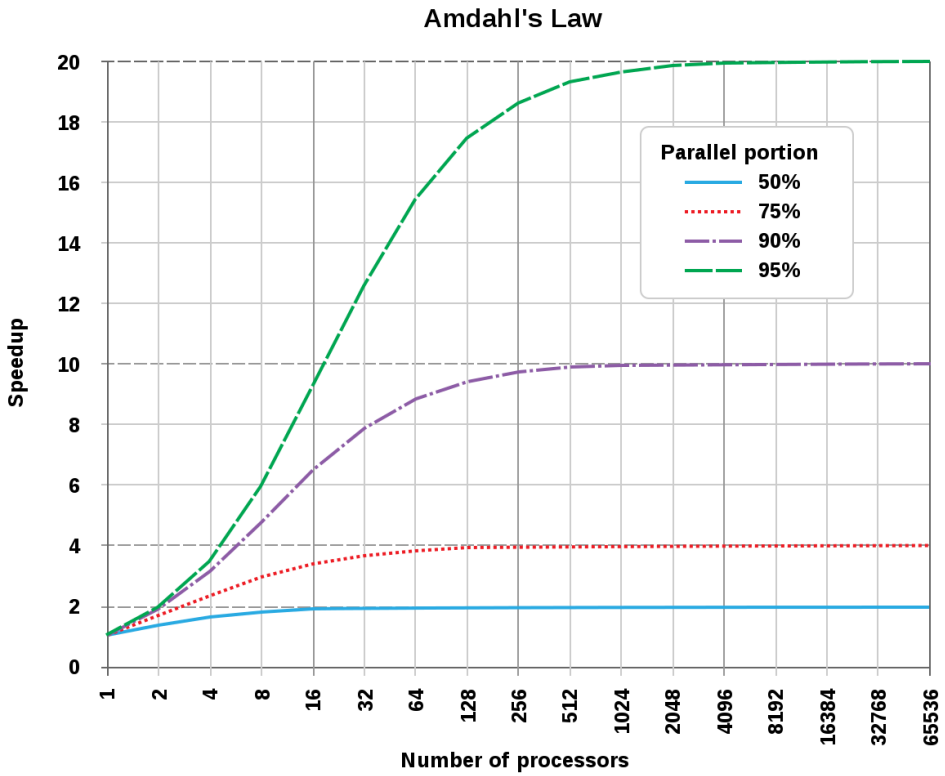

siliceum

Measure

The most important part!

Performance testing

- Reproducible
- As deterministic as possible
- Choice of input data is important
  - Size must be realistic
  - Represents the system in production or at its limits

With a single measurement

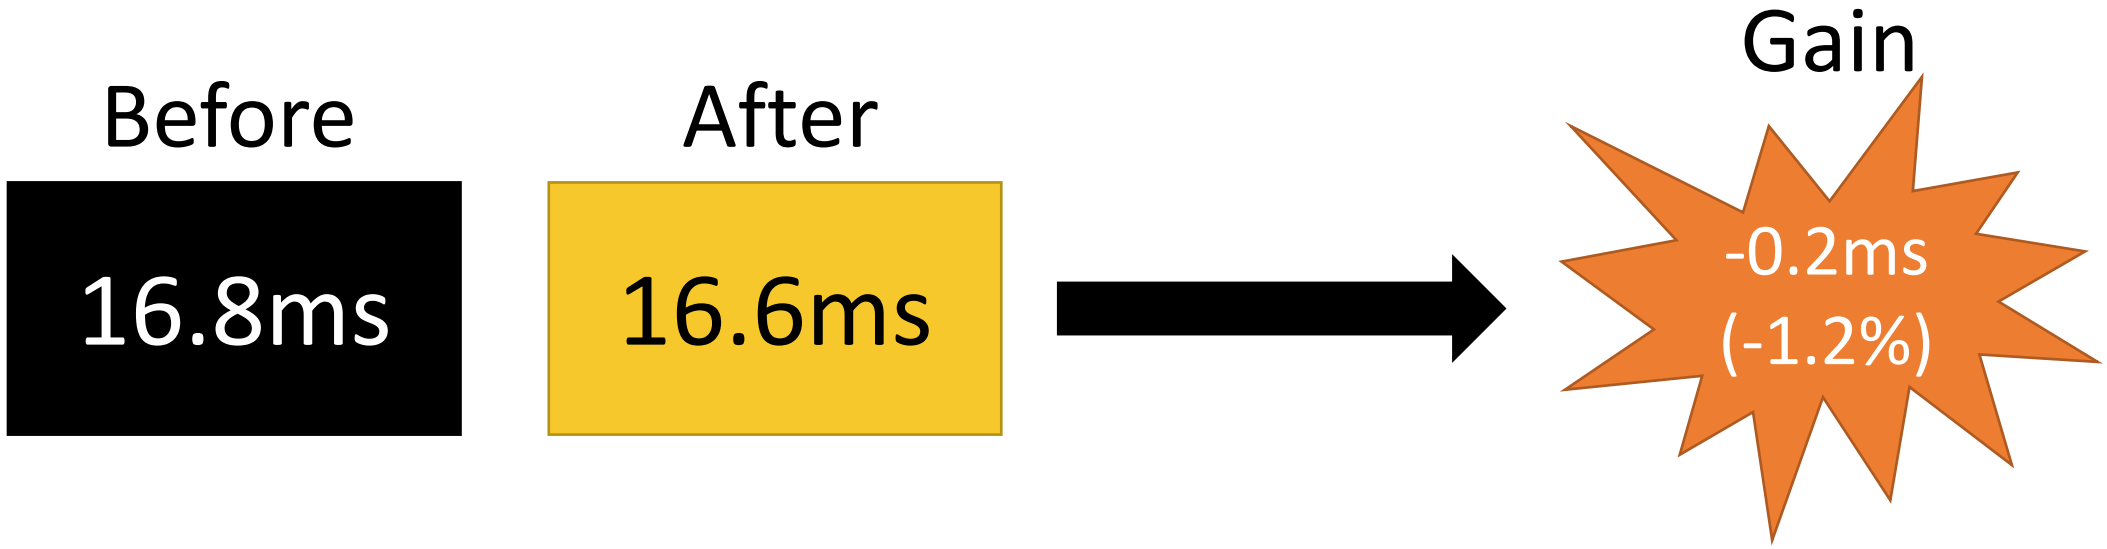

Ship it?

Environment stability

(Do not leave YouTube running in the background)

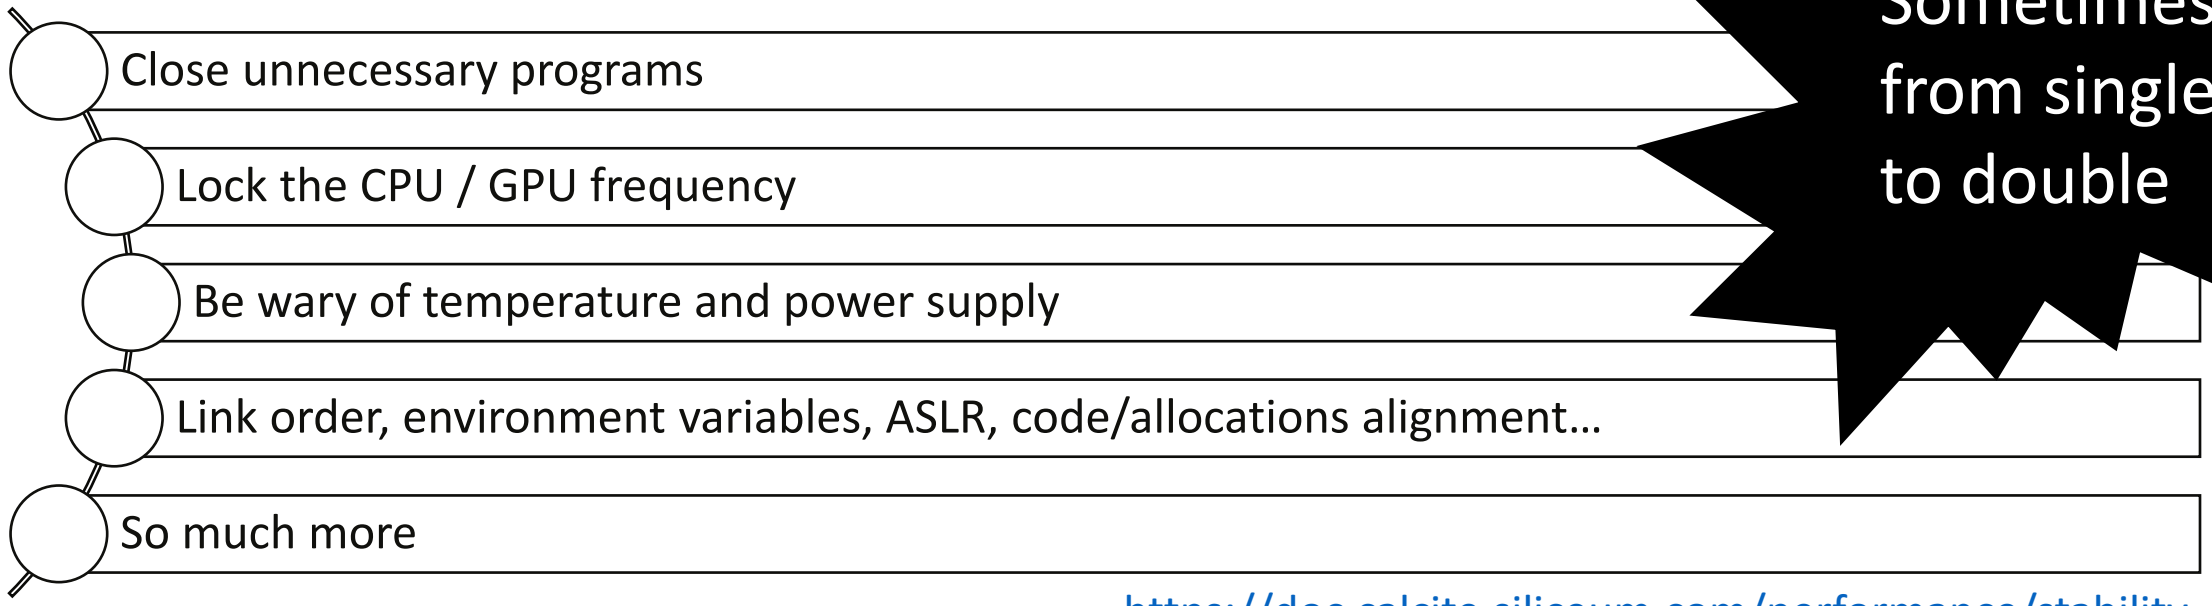

https://doc.calcite.siliceum.com/performance/stability

The solution

Repeat the experiment!

Use units correctly

- % only if no other way to describe it correctly
  - Cache miss rate ✅
  - Branch misprediction ✅
  - Timings ❌
  - Memory ❌
- No unit expressed in « ops / X » (FPS, throughput)
  - Ok for a marketing effect
  - Bad when looking at improvements
  - Same modifications in different order => different results
- NO AVERAGE → Use median / percentiles / statistical tests…

siliceum

Profilers

Tools of the trade

Sampling-based profilers

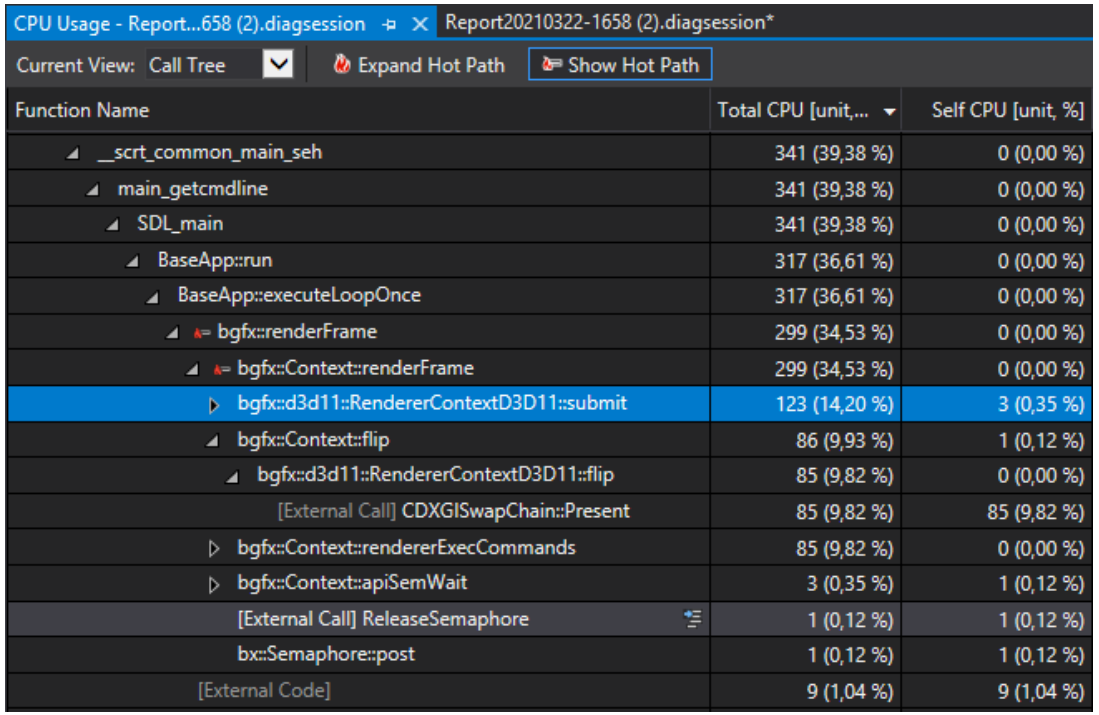

- Quick results ✓
- I/O not always sampled ⚠
- Variable precision ✗
- Time spent in function ≠Useful time in function ✗
- Linux Perf, Visual Studio, Intel VTune

Automatic instrumentation

(Avoid it 99,999% of the time)

```
Flat profile:

Each sample counts as 0.01 seconds.
  %   cumulative   self              self     total
 time   seconds   seconds    calls  ms/call  ms/call  name
 33.34      0.02     0.02     7208     0.00     0.00  open
 16.67      0.03     0.01      244     0.04     0.12  offtime
 16.67      0.04     0.01        8     1.25     1.25  memccpy
 16.67      0.05     0.01        7     1.43     1.43  write
 16.67      0.06     0.01                             mcount
  0.00      0.06     0.00      236     0.00     0.00  tzset
  0.00      0.06     0.00      192     0.00     0.00  tolower
  0.00      0.06     0.00       47     0.00     0.00  strlen
  0.00      0.06     0.00       45     0.00     0.00  strchr
  0.00      0.06     0.00        1     0.00    50.00  main
  0.00      0.06     0.00        1     0.00     0.00  memcpy
  0.00      0.06     0.00        1     0.00    10.11  print
  0.00      0.06     0.00        1     0.00     0.00  profil
  0.00      0.06     0.00        1     0.00    50.00  report
```

- Deceptive overhead ✗
- « Useful » to know the number of function calls
- Gprof, tyoma/micro-profiler, …
- Read this StackOverflow

Manual instrumentation

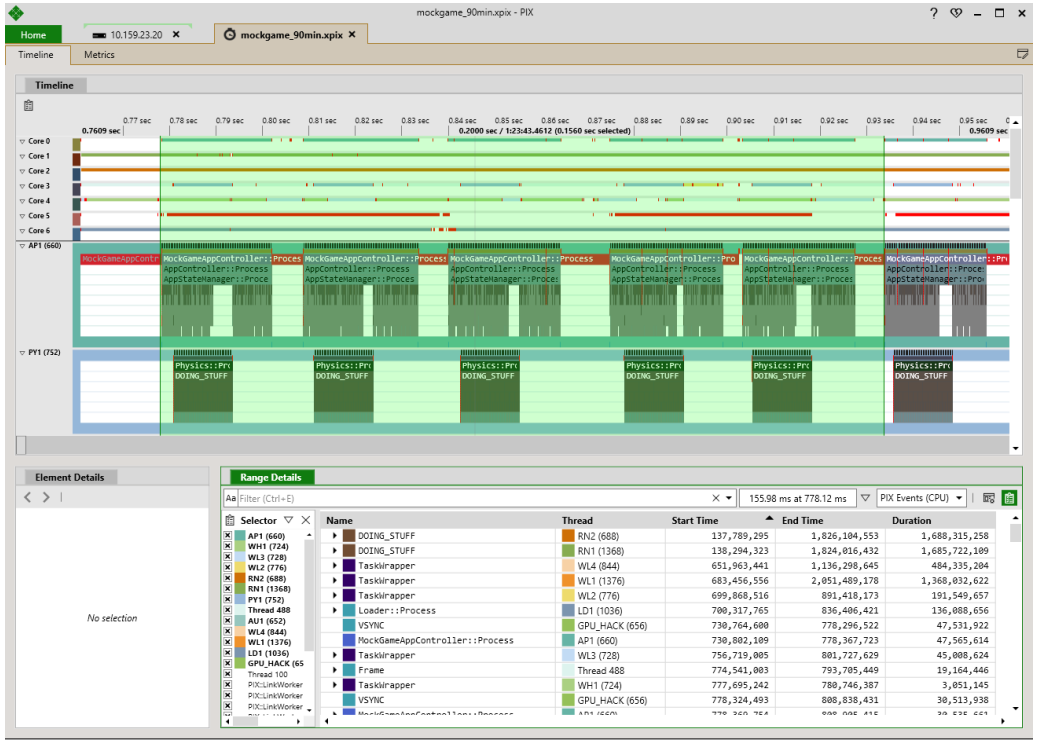

- Precise! ✓
- Macro view! ✓
- Business logic hierarchy ✓
- No details ✗
- Requires code changes ✗
- Remotery, easy_profiler, …

Instrumentation + Sampling

Best of both worlds

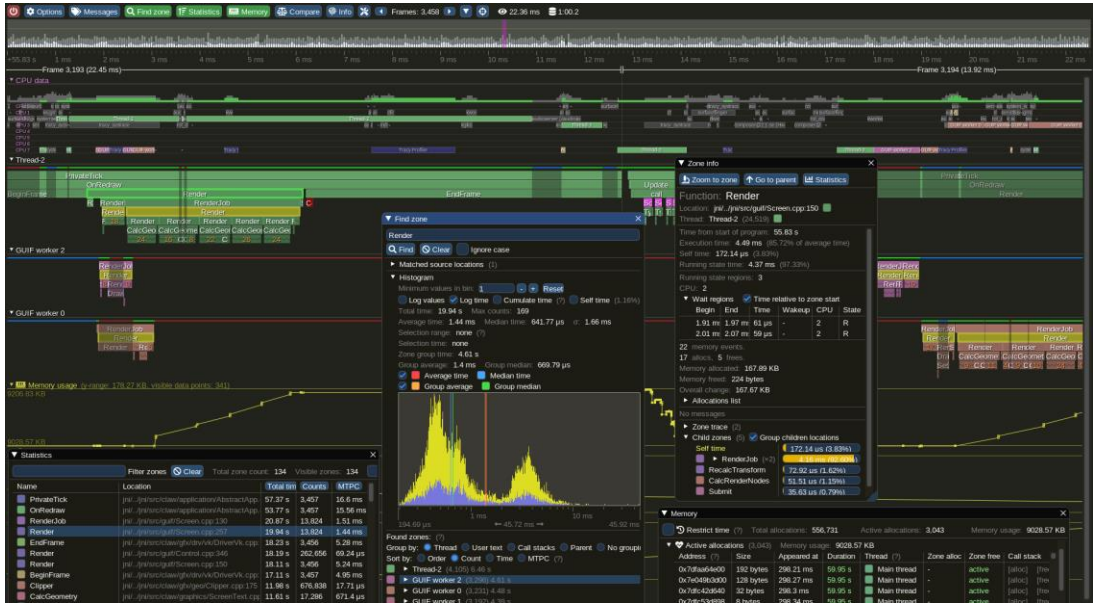

- Microsoft PIX for Windows
- Tracy
- Optick
- Superluminal
- Perfetto
- …

Profiler wishlist

- ❑Lightweight
- ❑Accurate
- ❑Multi-threading support
  - ❑Thread naming
  - ❑Context switch details
- ❑I/O measurements
- ❑Memory measurements
- ❑Custom data sources
- ❑Cross-platform
- ❑A good UI/UX
- ❑Can save sessions
- ❑Fibers/coroutines support?
- ❑Can be attached to a running app
- ❑CI support?
- ❑…

siliceum

Macro-optimization

Looking at the big picture

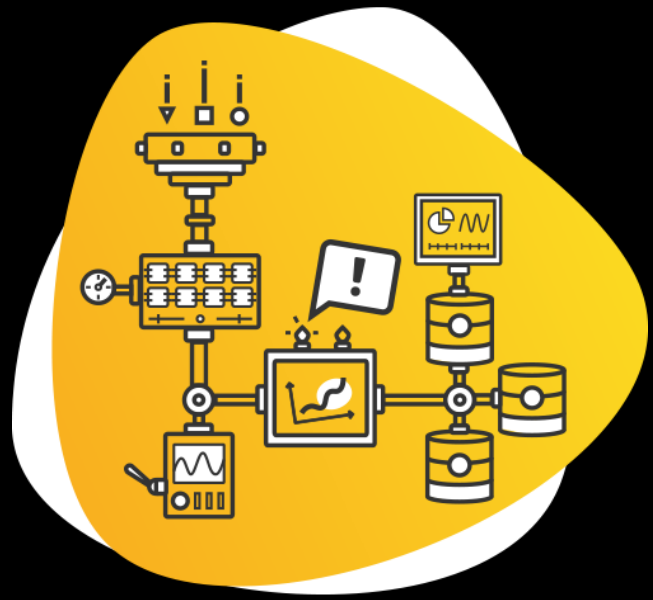

Compilation flags

Quick to check and test

GCC/Clang

- -O3
- -march=native
- -flto
- -ffast-math ⚠

MSVC

- /O2
- /arch:???
- /GL, /LTCG
- /fp:fast ⚠

Do nothing

The best of optimizations

- Avoid unneeded processing
- Avoid unneeded initialization/copies
- Avoid unneeded (re)allocations (use `reserve()`!)
- Avoid using mutexes (and reduce the code it guards)
- Avoid duplicated function calls (abstractions!)

Adapt your data or design

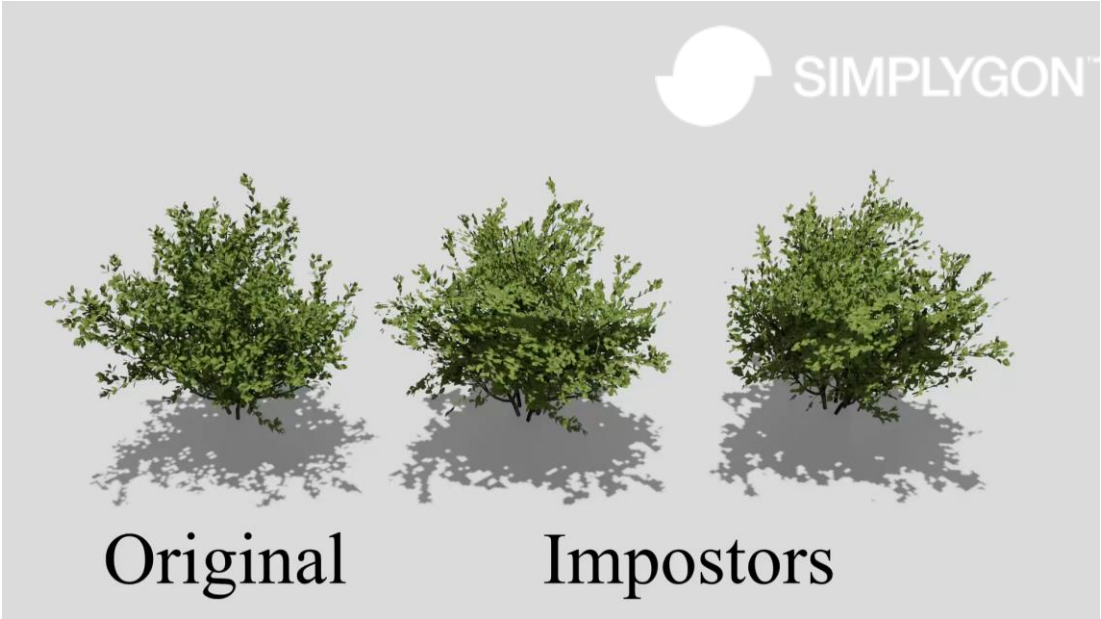

Adapt your data or design

Yann Richard ©

In order of efficiency:

- Do less
- Do it faster
- Prepare data (compress, pre-compute, …)
- Multi-core (parallelism, latency hiding, …)
- Caching (do this last or you'll regret it!)

siliceum

War story: WebKit GC in a game

WebKit and real-time

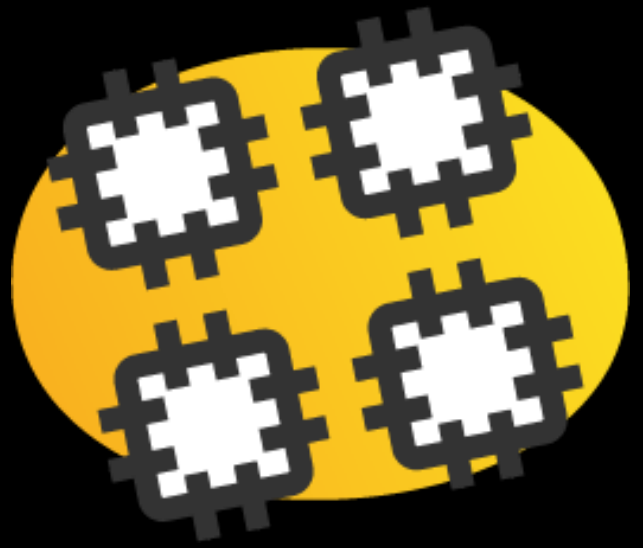

Some context

- WebKit used to display dynamic UIs, game components…
- Target for a game is 60|30fps => 33|16ms
- No JIT on console
- « Old » WebKit version from ~2016
- Frequent freezes could be observed

Remember:

Measure 1st

(we already have our target: 60/30fps)

Start collecting data

Conquered land

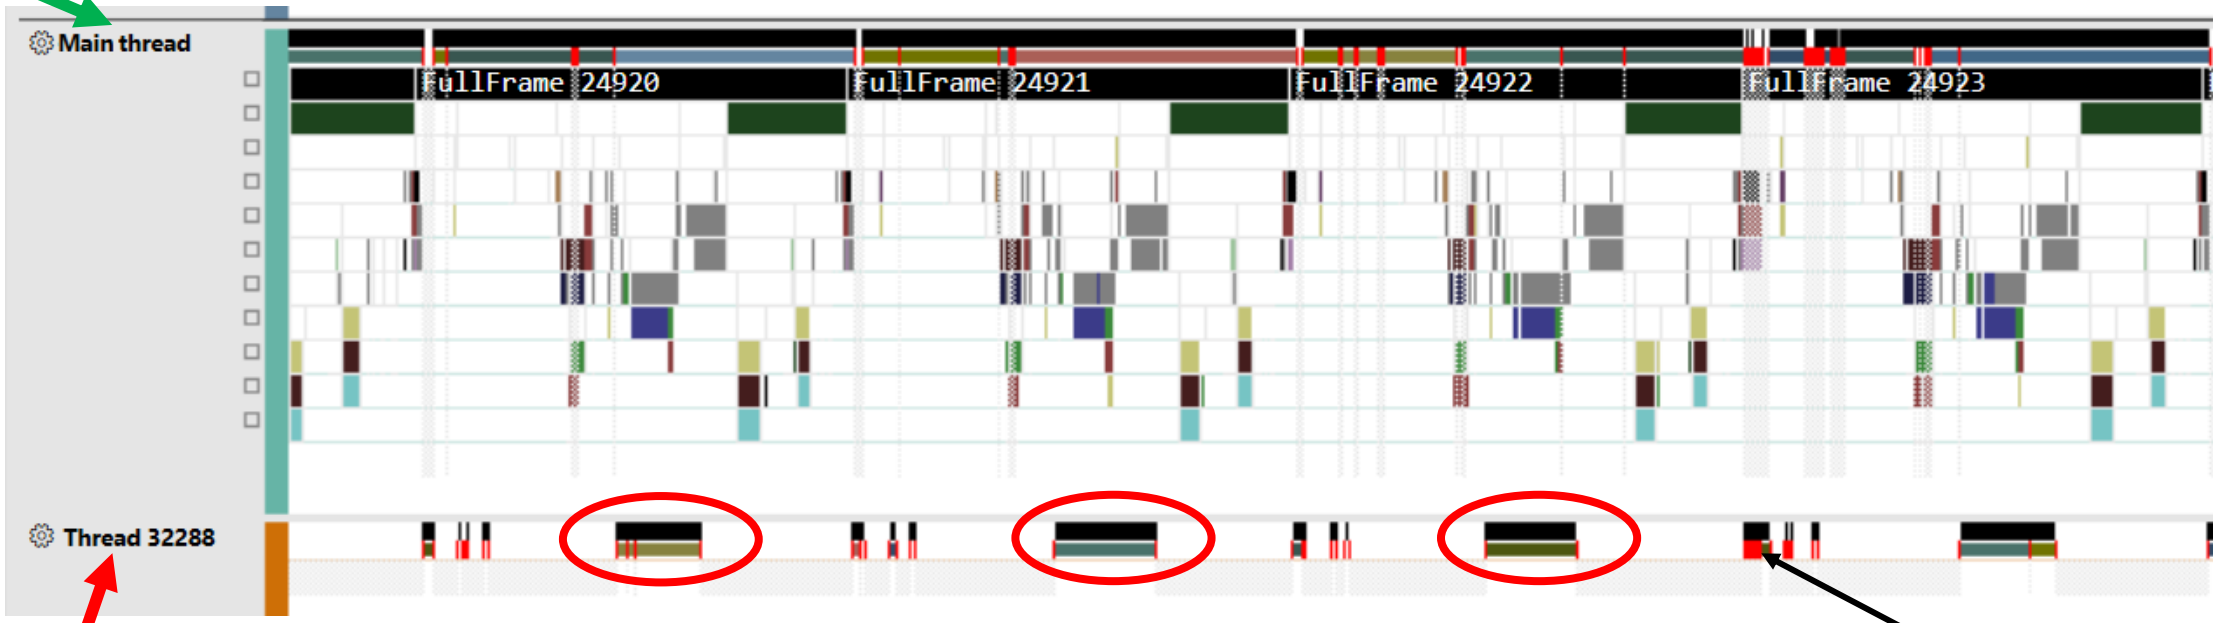

Red flag

No instrumentation, only sampling & context switches

What it looks like in the code

```cpp
//////////////////////////////////
bool Engine::Update()
{
    PERF_REGIONF();

    UpdateInput();

    UpdateMessages();
```

```cpp
//////////////////////////////////////////////////////////////////////
void Engine::UpdateLogic()
{
    PERF_REGIONC("UpdateLogic", PerfMarkers::Category::GameLogic);
    PERF_TAG("PlayerName", name);

    // Update unit

    PERF_TAG("Position", 123.0f, 456.0f, 789.0f);
    PERF_TAG("Health", 100);
    PERF_TAG("Address", (uint64_t)this);
```

The garbage collector monster

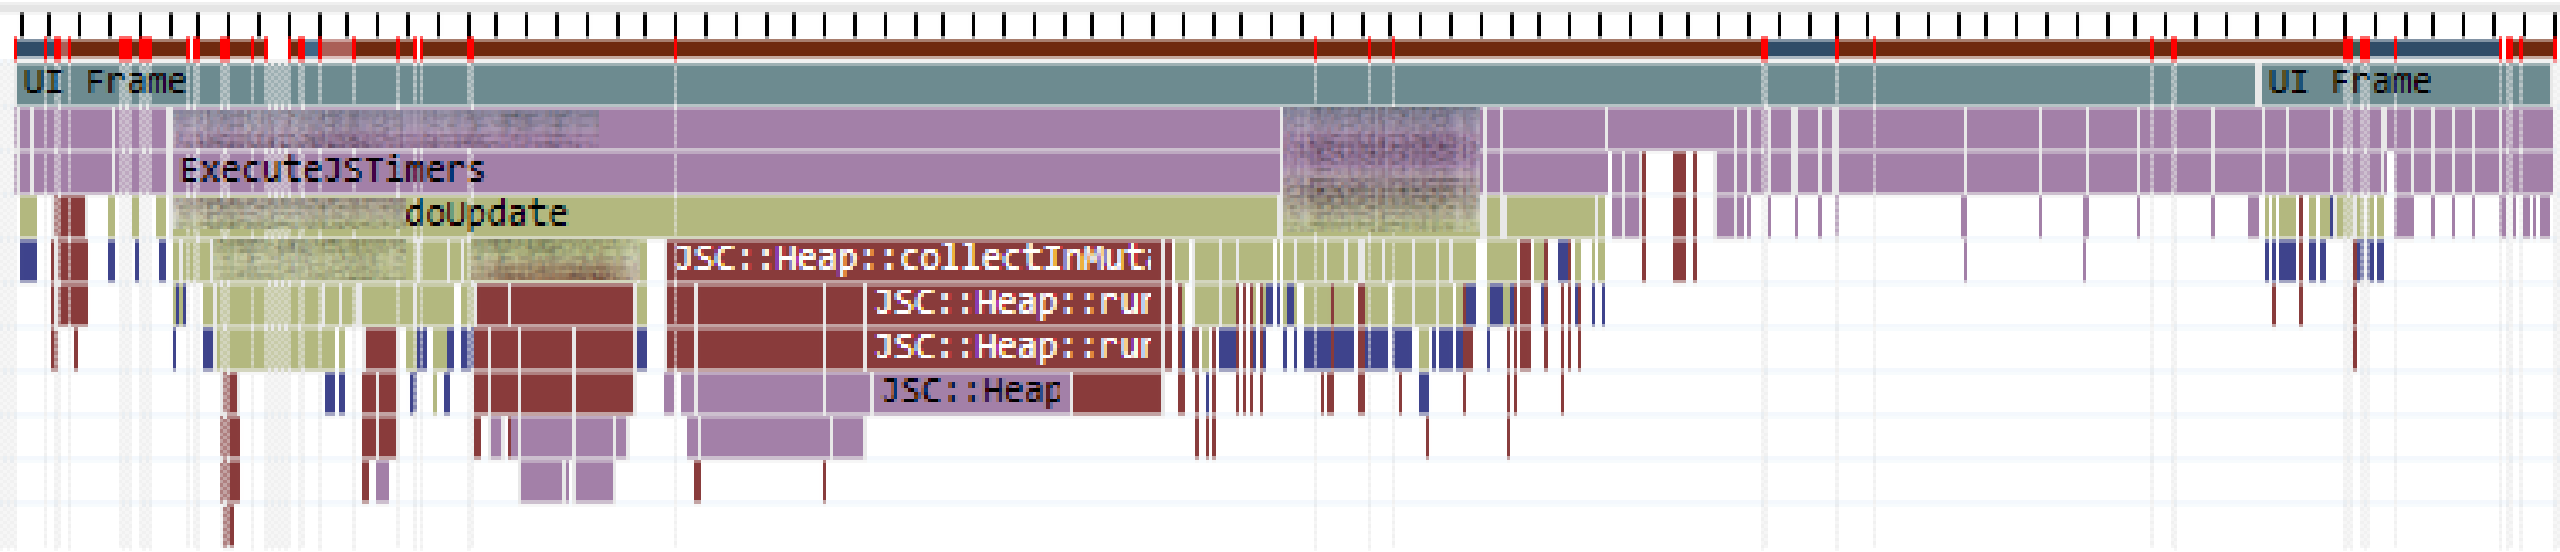

What do we do now?

Many options

- ❑ Tweak the garbage collector parameters
- ❑ Optimize the garbage collector
- ❑ Make it less blocking
- ❑ Rewrite the code in C++
- ❑ Reduce the amount of data to process
- ❑ ~~Ship it~~

Reduce the amount of data to process

Big wins without diving into WebKit

- Analyze what generates memory usage
- Fix JS code to generate less memory allocations
- Sadly not enough
  - Still no control over when the GC will happen
  - Still have frequent « freezes » due to GC

Eden vs Full GC

- Sometimes GC takes 70ms, sometimes « only » 10ms?
- Full GC
  - Iterates over all allocations!
- Eden GC
  - Iterates only over objects allocated/changed since the previous GC
  - This is what we want most of the time in a game loop

Avoiding full GC

- Full GC always triggered after eden
- Solution:
  Trigger only if size after GC is
  - above full GC threshold
  - X% greater than previous post-full GC size

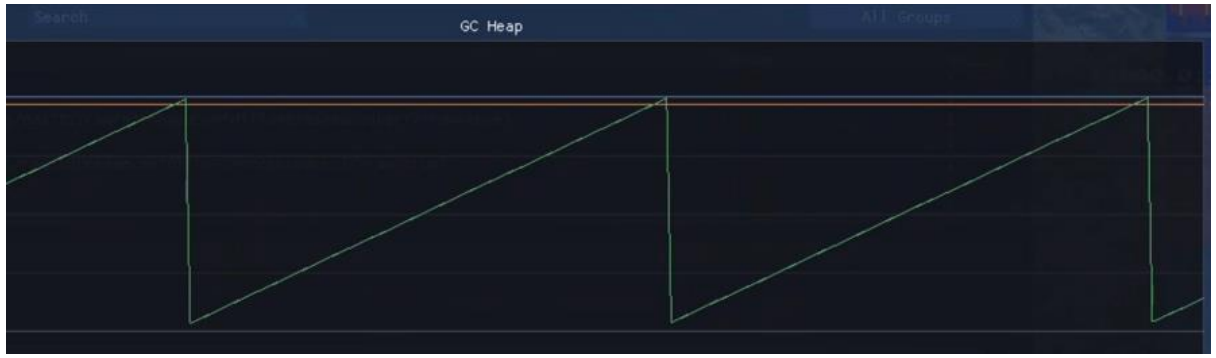

Eden GC size threshold

Current memory size

GC may pause execution to run

Scheduler not a good fit for realtime purposes

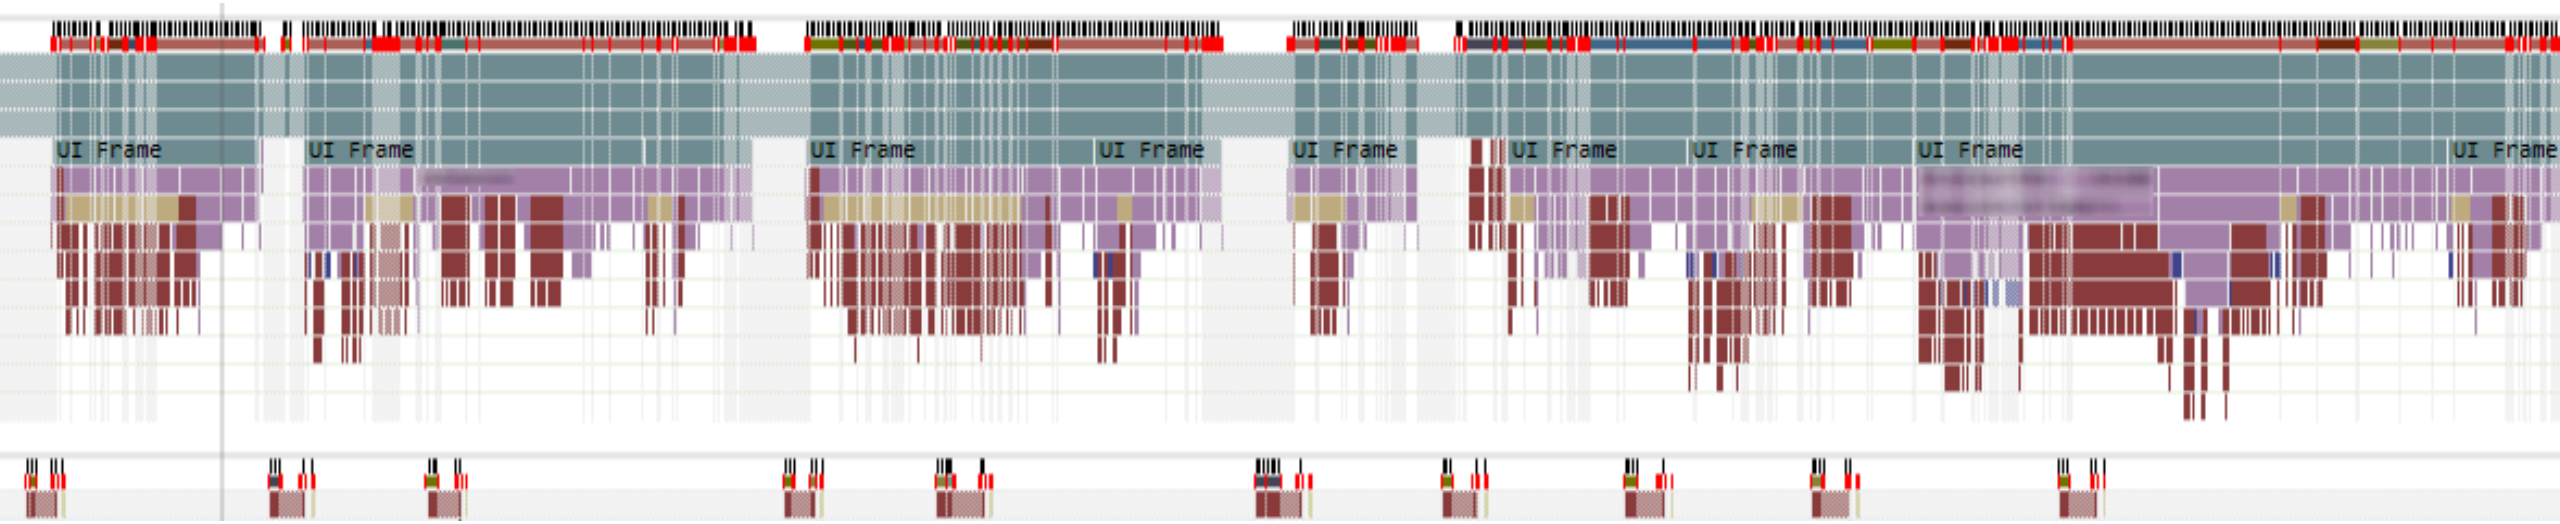

Prologue & epilogue of GC passes are expensive

stop/resumeThePeriphery

Actual useful work

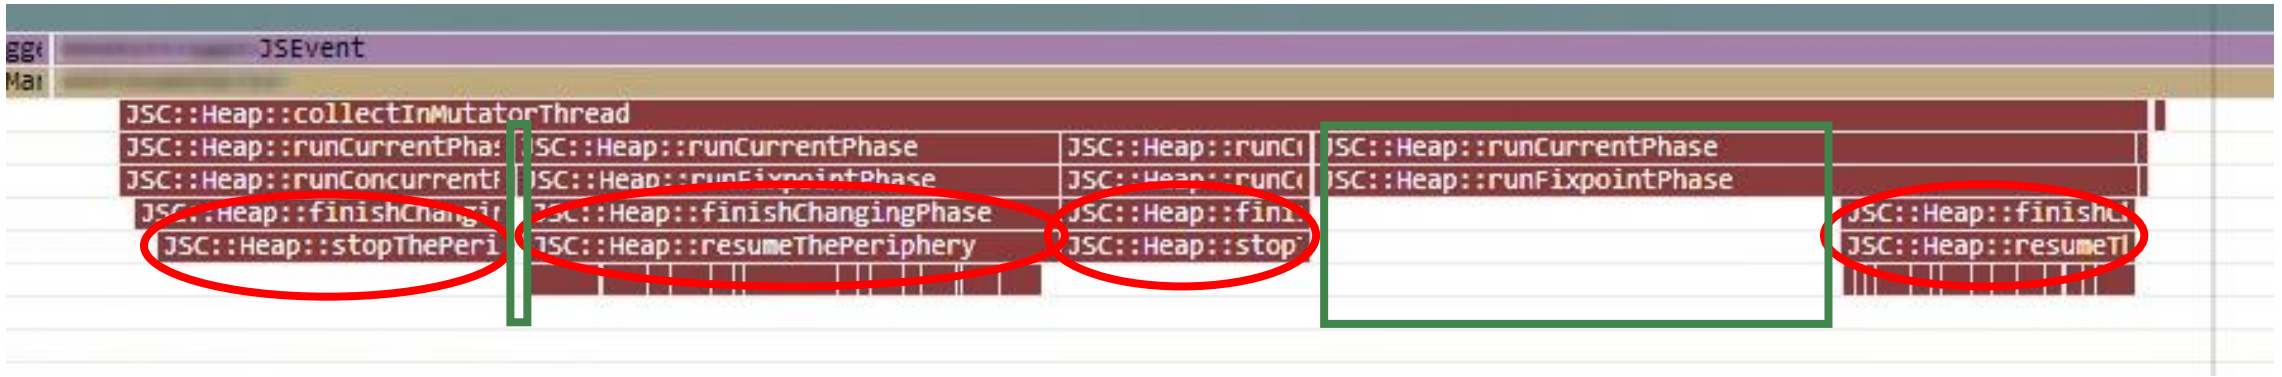

A normal frame

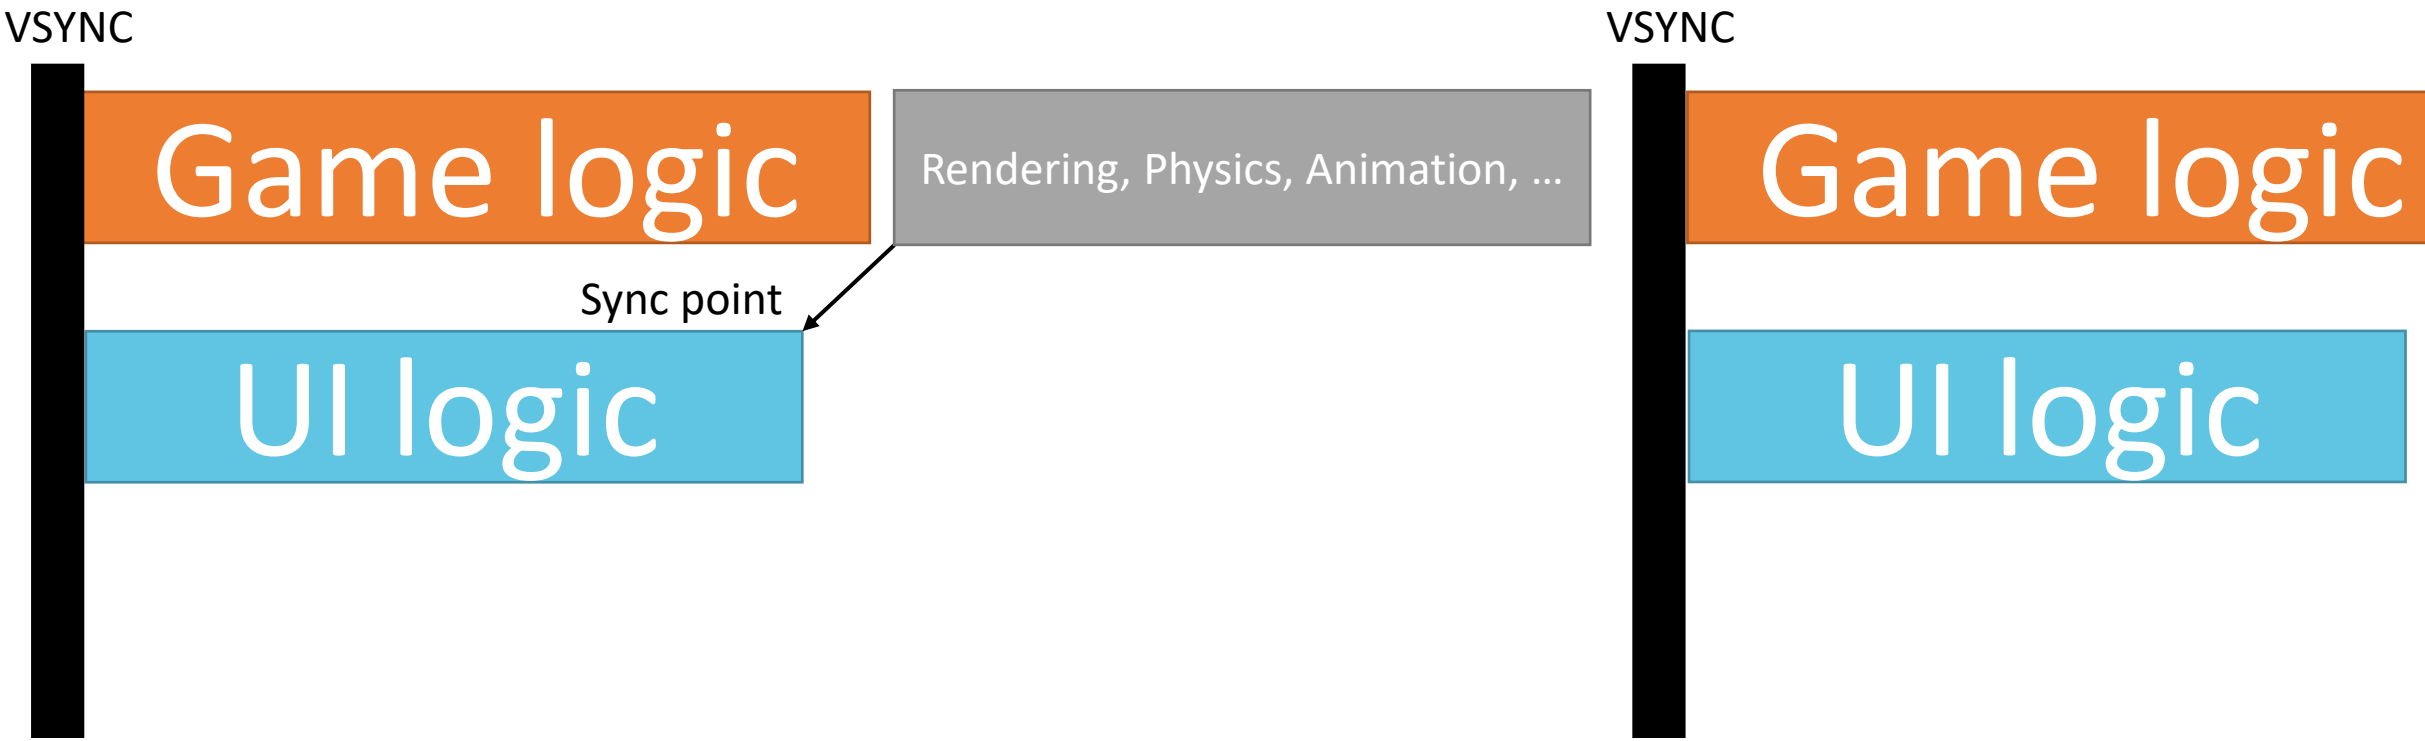

Here be dragons (GC)

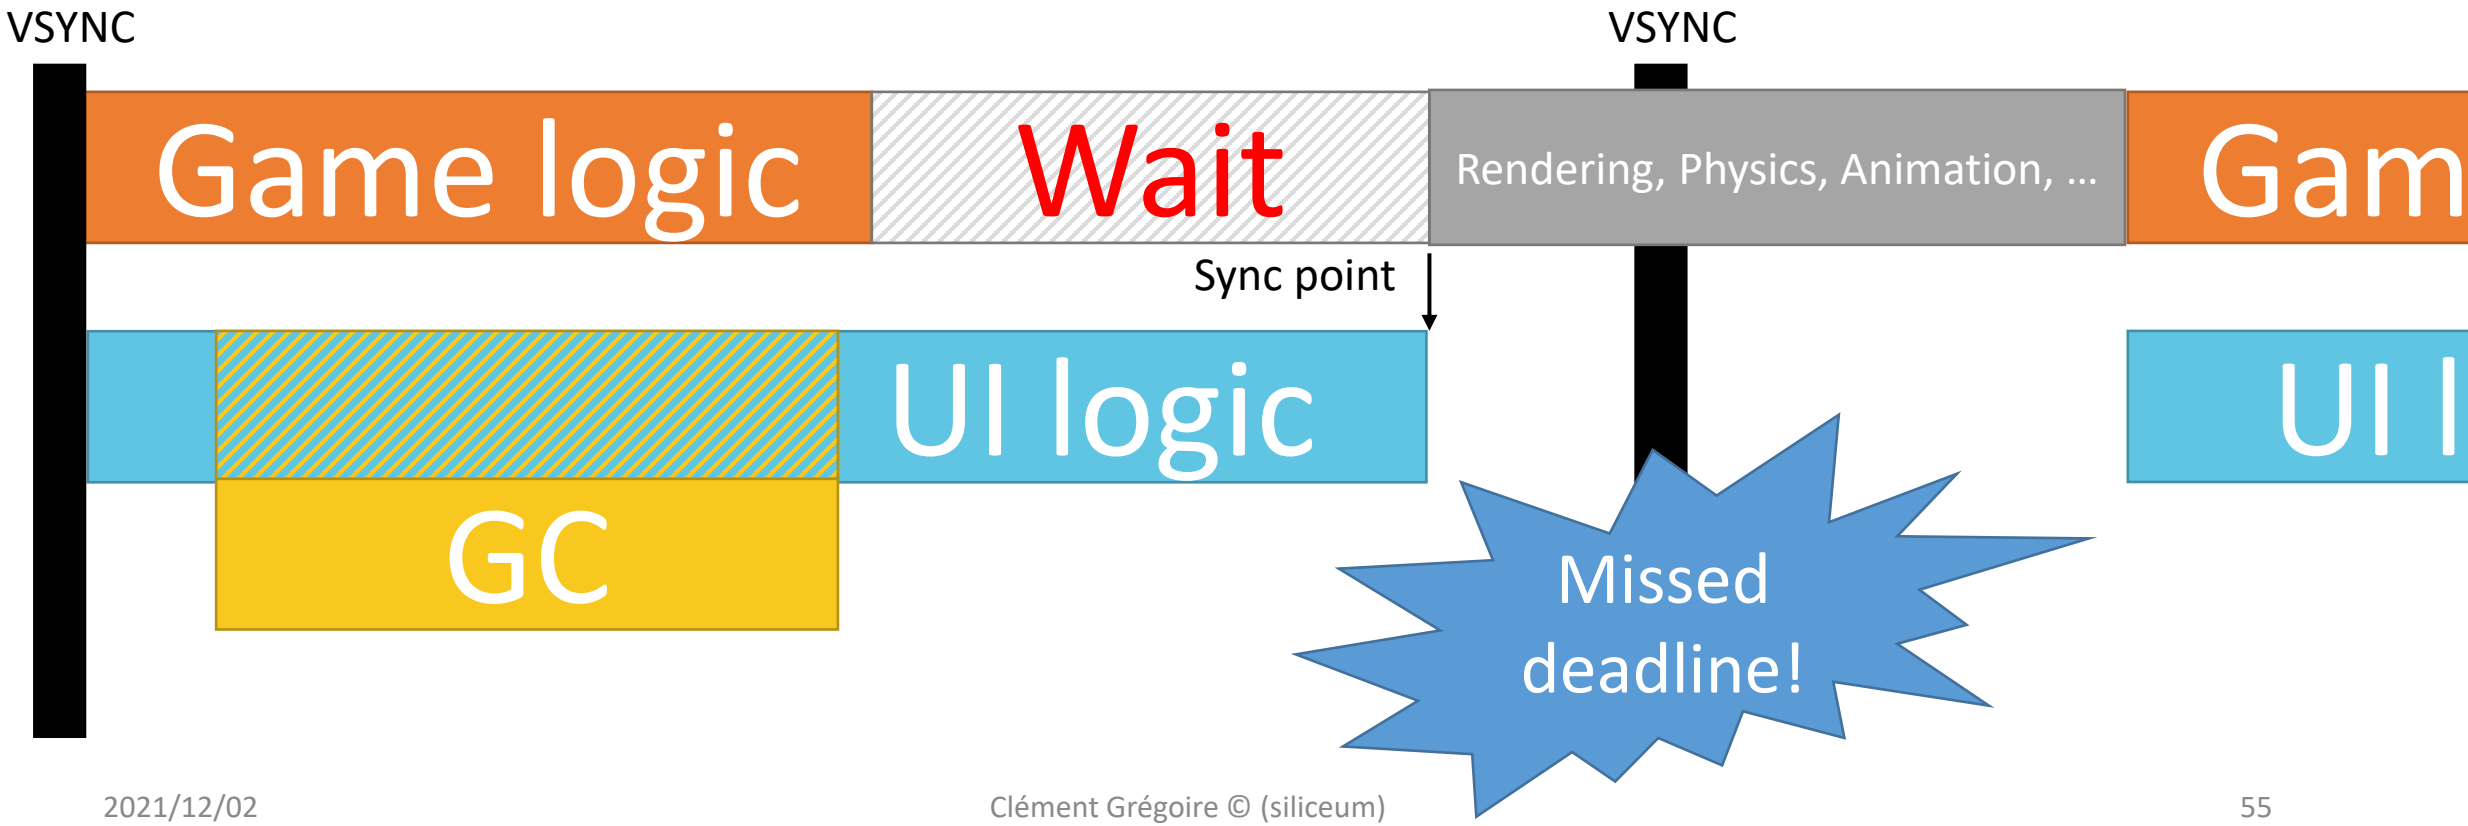

The solution

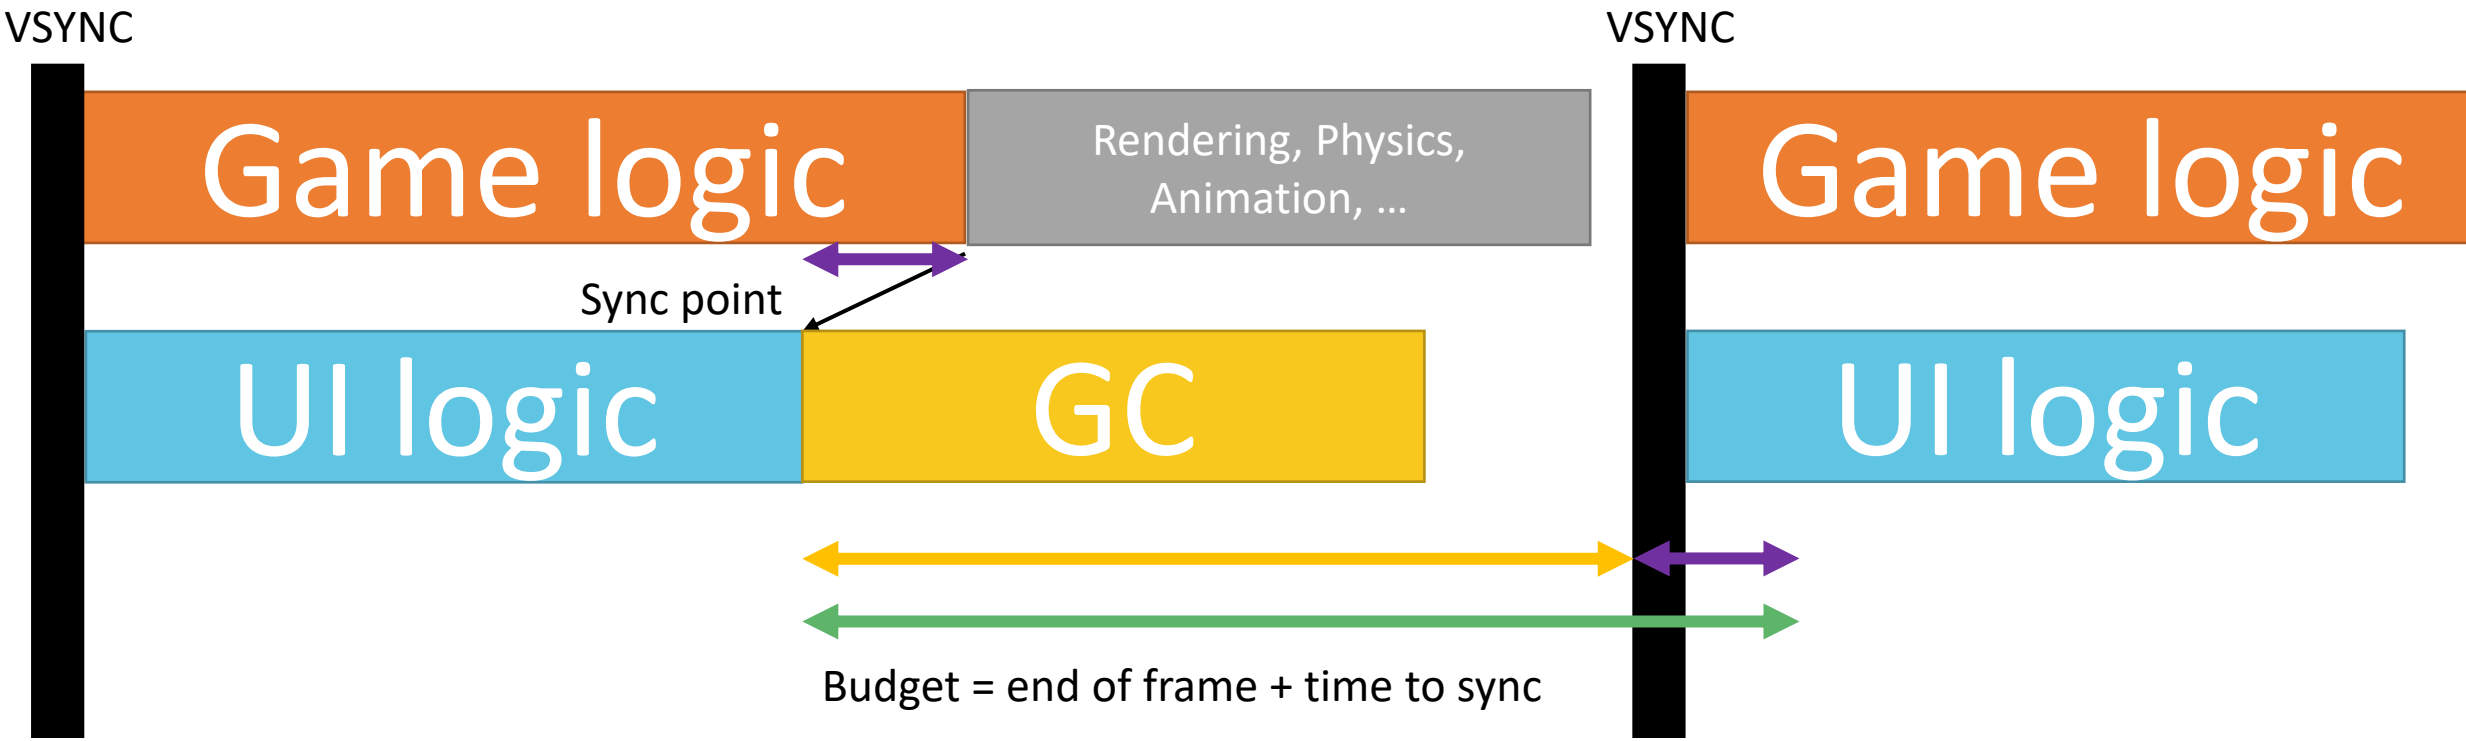

Bonus point


JSEvent
JSC::Heap::runCurrentPhase
JSC::Heap::runFixpointPhase

Update JSCore to benefit from multi-threading

Take advantage of all the cores!

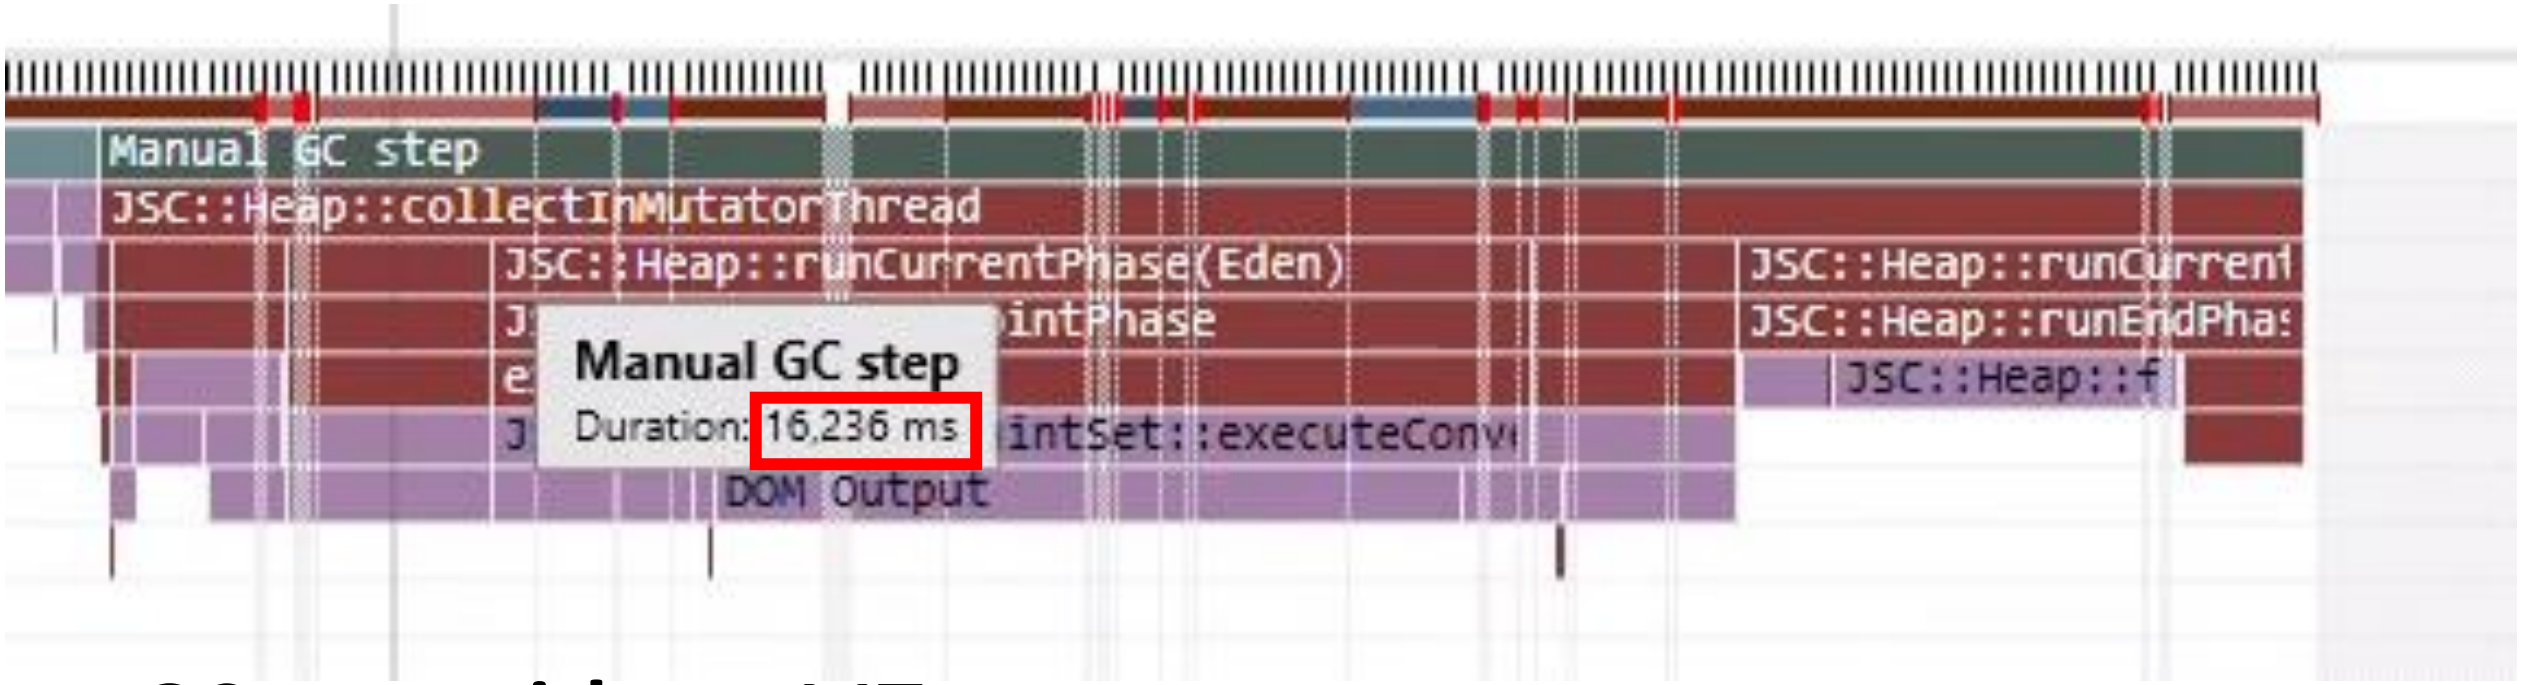

GC step **without** MT

Update JSCore to benefit from MT?

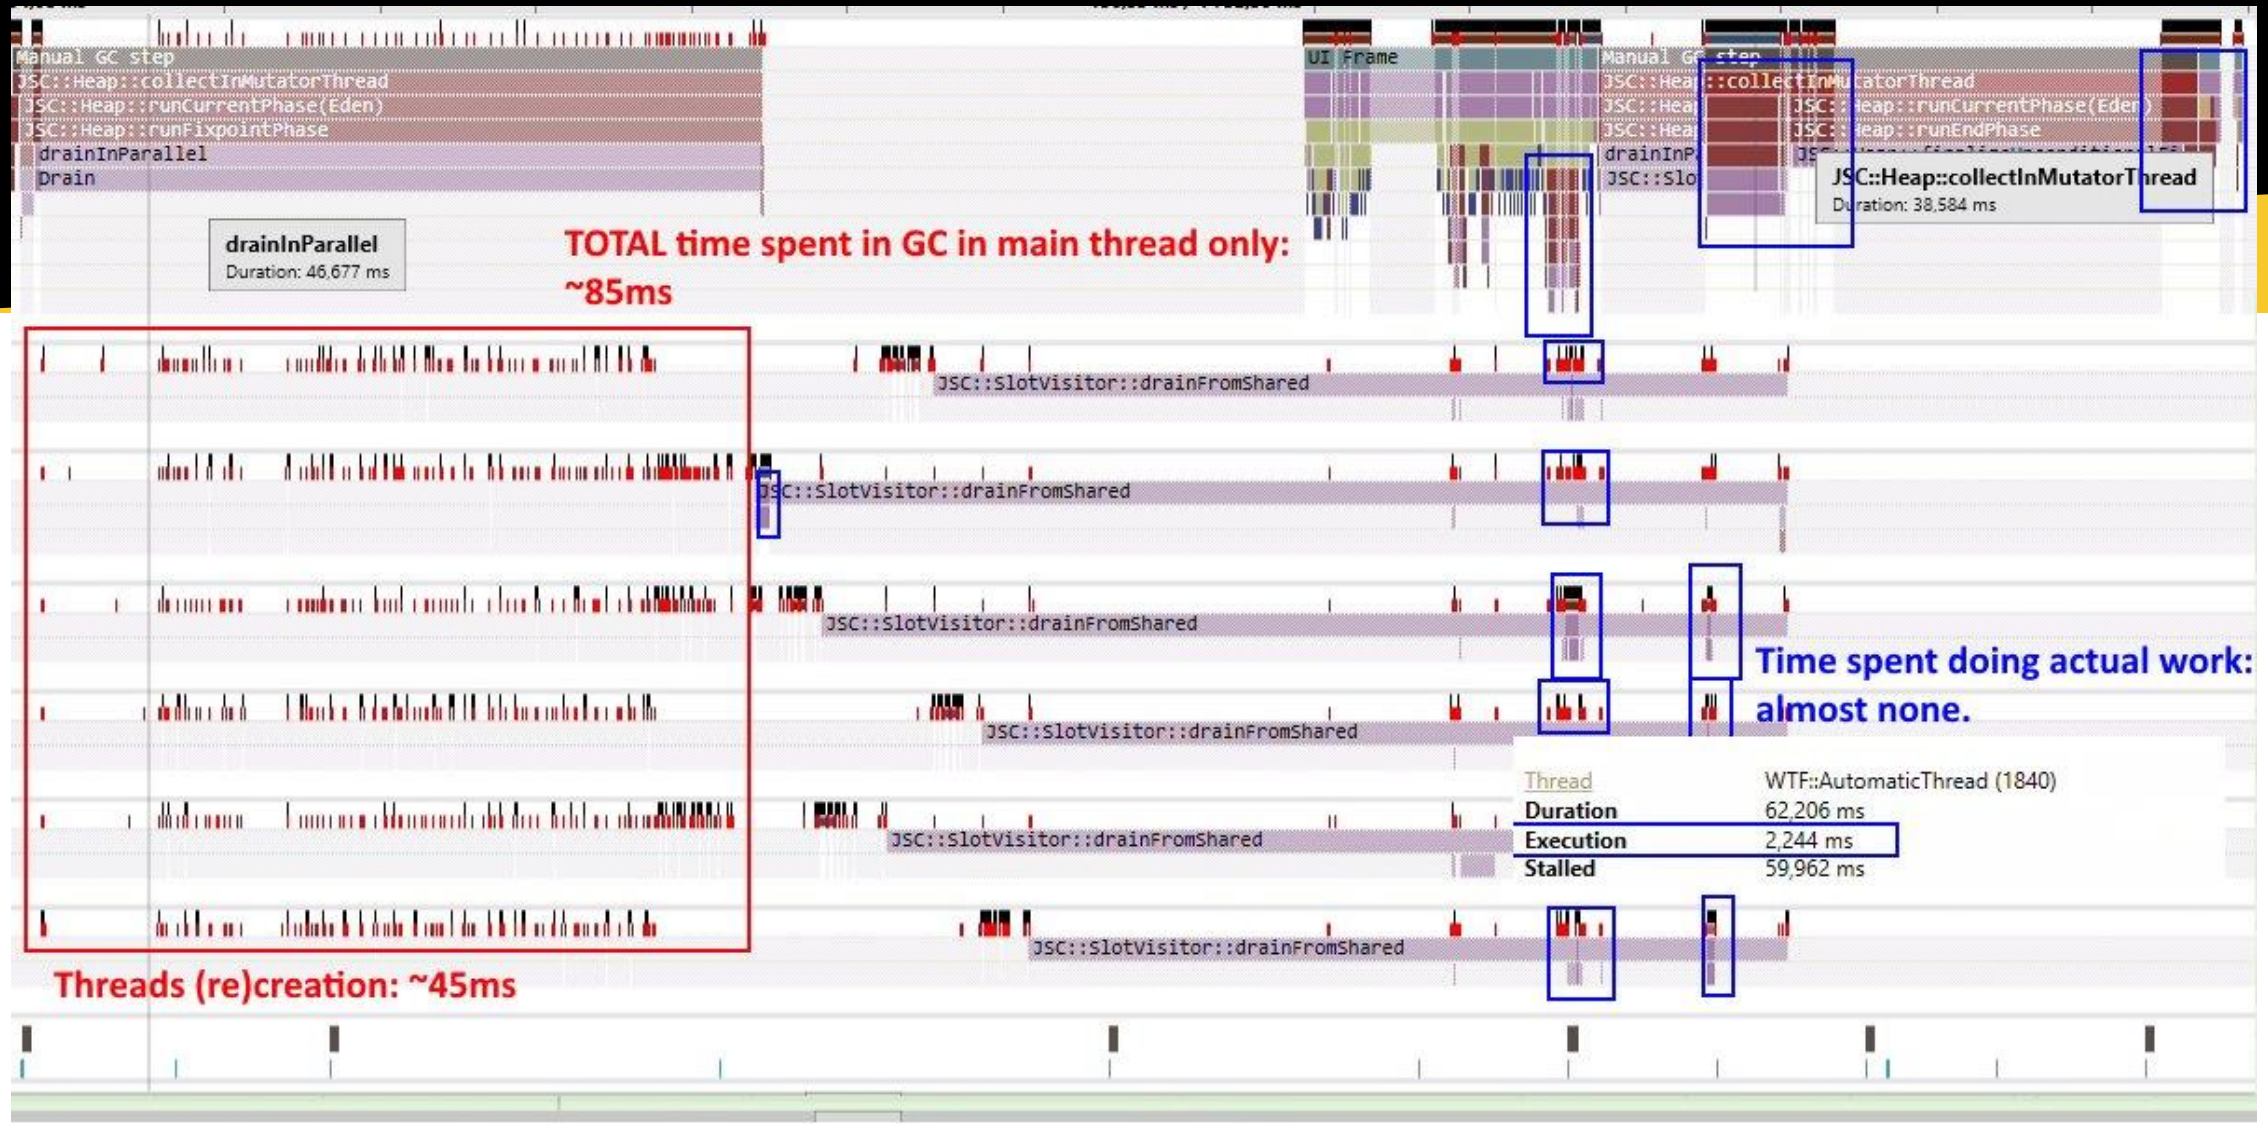

Update JavaScriptCore to benefit from MT

Does not scale as-is

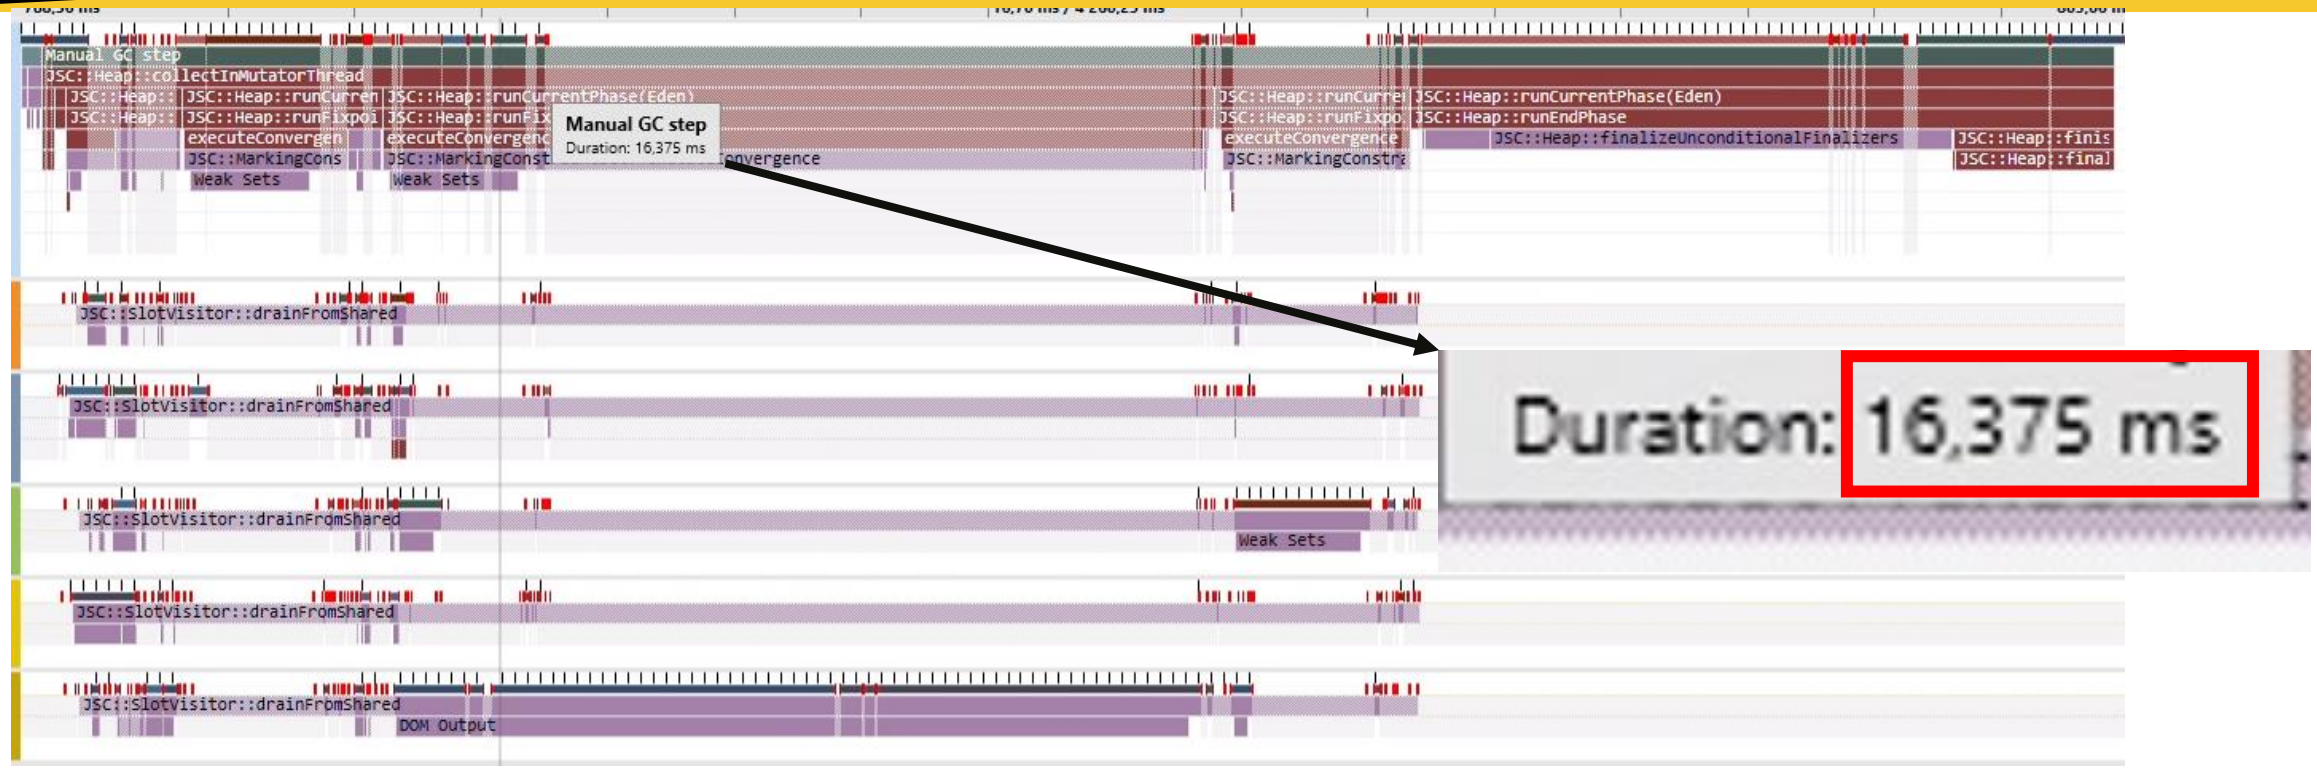

~~Update JSCore to benefit from multi-threading~~

The wild goose chase

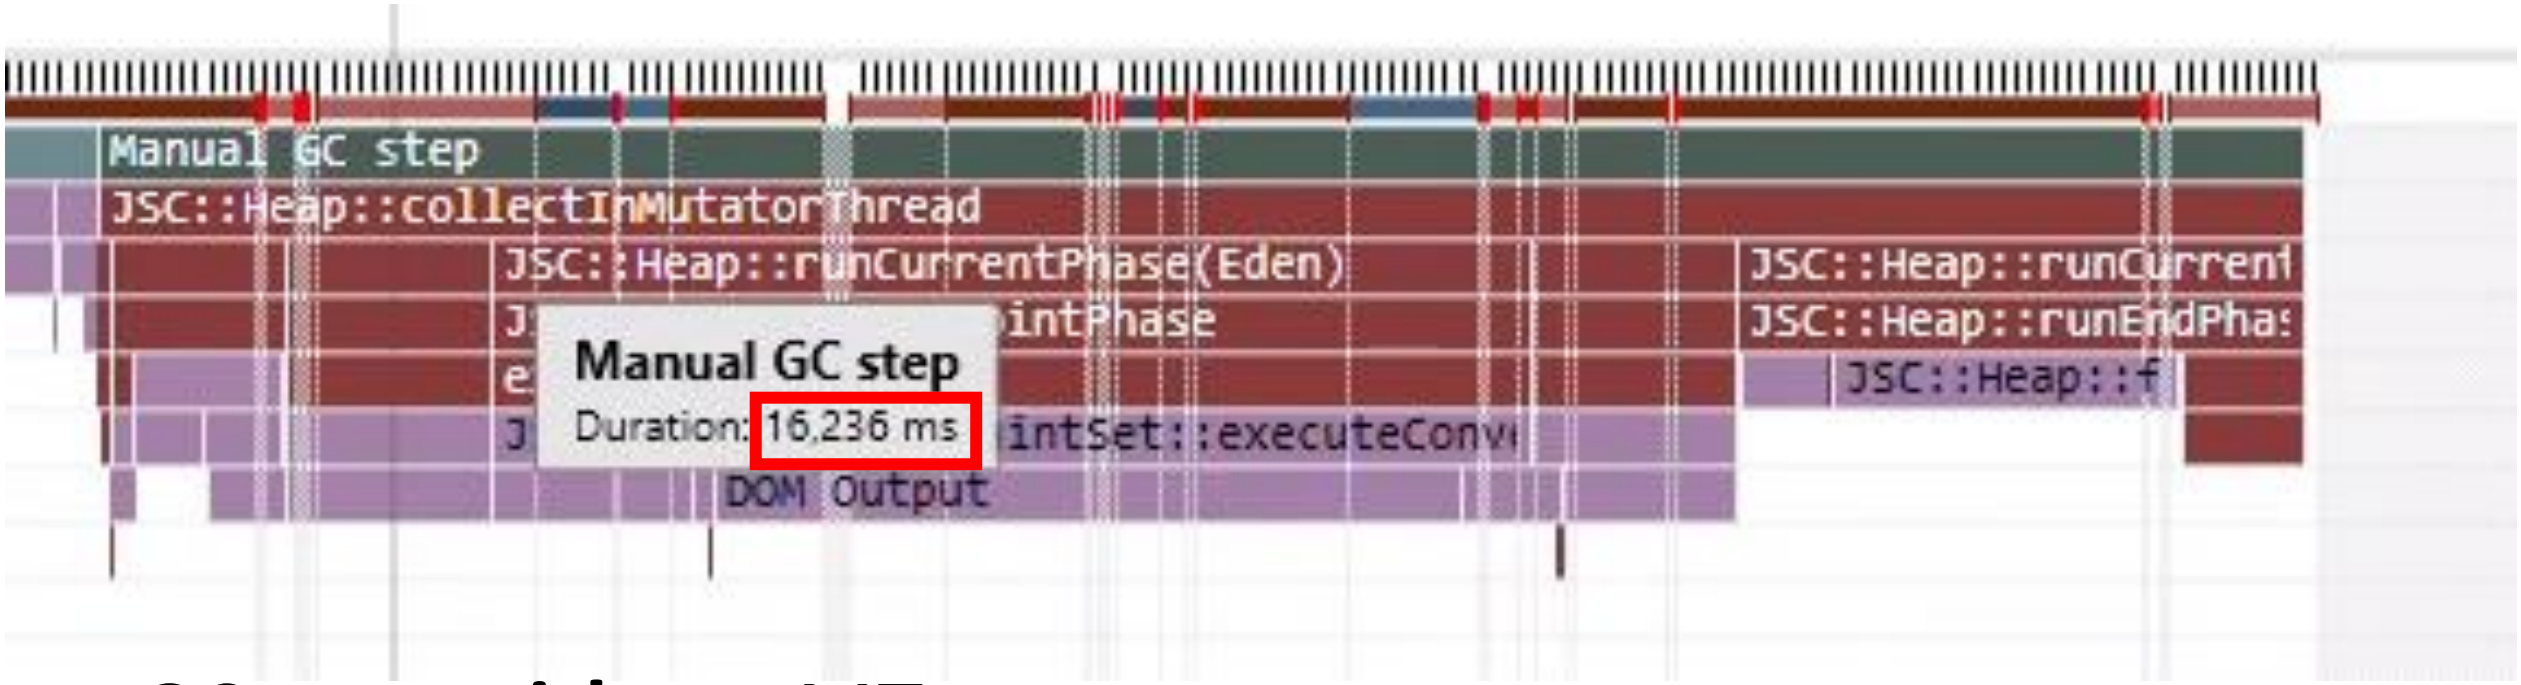

GC step **without** MT

Many more optimizations

Not really C++

- Reduce string copies for C++ ⇔ .js communications
- Reduced allocations and indirections in WebKit
- No getters/setters => slow path in VM
- Reduce DOM interactions
- Functions forcing layout (Element.getBoundingClientRect())

If you had to remember one thing

Measure, measure, measure!

(and understand what you measure)

Thank you!
(contact me)

Email:
clement@siliceum.com

Twitter: @lectem
Discord: #include

siliceum

Appendix

Latency Numbers Every Programmer Should Know

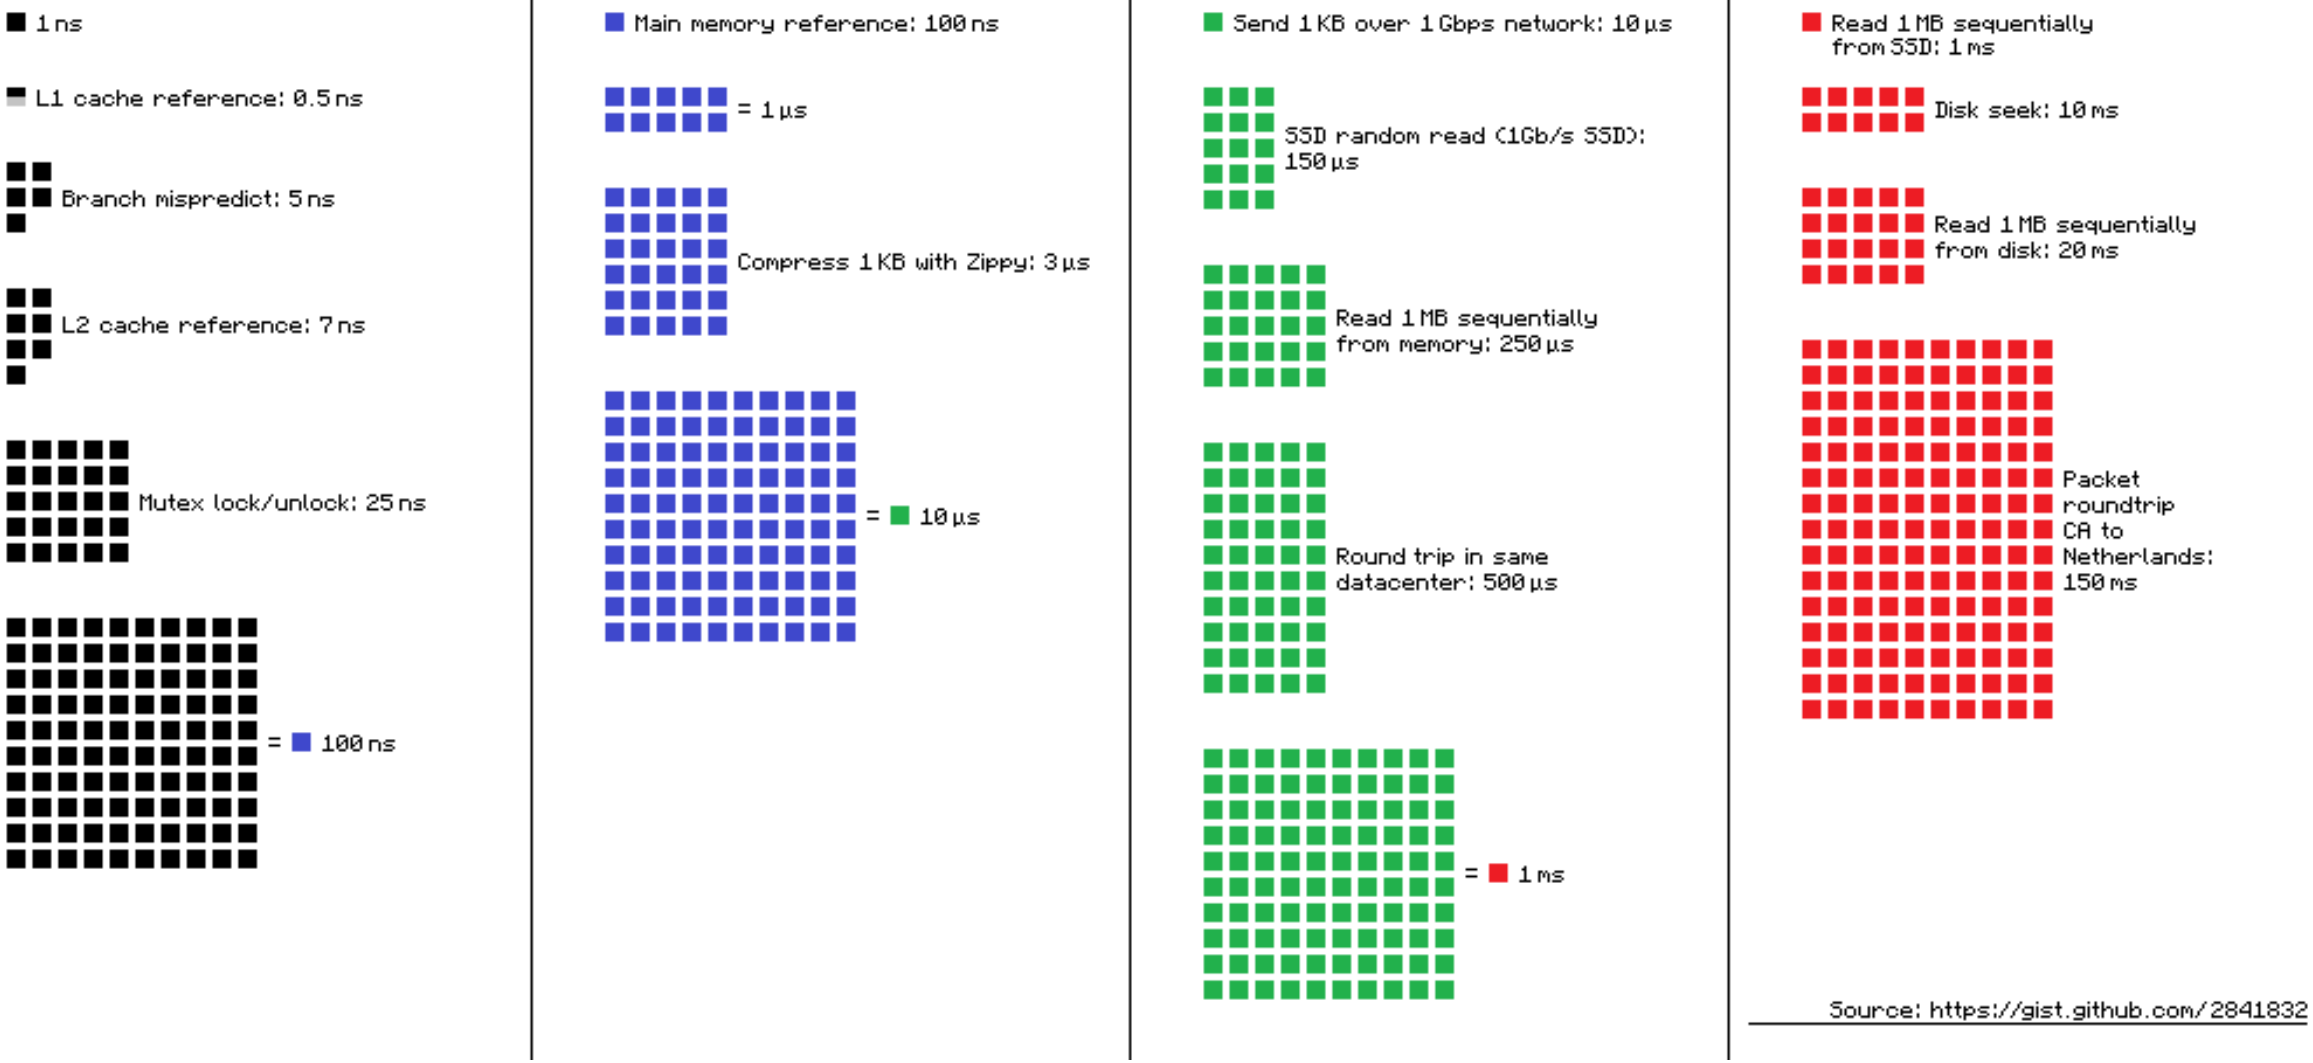

Not all CPU operations are created equal